

Esteban Sitek

Analytics Engineer

estebanbenjaminsitek74@gmail.com · +54 93885170292 · San Salvador de Jujuy, Jujuy, Argentina · www.linkedin.com/in/esteban-benjamin-sitek-346771398/

PROFESSIONAL SUMMARY

Senior Software Engineer with **12+** years across healthcare, fintech, and e-commerce, building scalable data pipelines and analytics infrastructure using **SQL**, **Python**, and cloud platforms. Experienced in translating complex business requirements into reliable enterprise data assets, implementing **Airflow** orchestration, and designing data models that drive strategic decision-making at scale.

SKILLS

- Languages:** Python, SQL, JavaScript, TypeScript, Java, HTML, CSS
- Data & ML:** Pandas, NumPy, PySpark, Spark, dbt, Airflow
- Database:** Snowflake, PostgreSQL, BigQuery, Redis, MongoDB, Elasticsearch
- Analytics & BI:** Tableau, Looker, Power BI, Metabase, Google Analytics
- Testing:** Jest, Pytest, Playwright, Cypress, k6
- Cloud & DevOps:** AWS, Google Cloud, Docker, Kubernetes, Terraform, GitHub Actions, CloudWatch, Datadog, Grafana, PagerDuty
- Others:** Git, GitHub, RESTful API, GraphQL, Agile, Jira, Confluence, Data Governance, Data Lineage, Star Schema, Alation, Collibra, CI/CD

EXPERIENCE

Publicis Sapient Jun 2024–Present
Senior Software Engineer

- Led the design and implementation of enterprise data pipelines for healthcare analytics, building foundational data assets using **Python**, **SQL**, and **Snowflake** to serve as the single source of truth across multiple business domains.
- Collaborated with healthcare stakeholders and business analysts to translate complex clinical and claims data requirements into scalable data models, establishing clear data lineage and governance frameworks that reduced analytics rework by 30%.
- Architected end-to-end data orchestration workflows using **Airflow** DAGs, scheduling 40+ daily data jobs that process healthcare claims and population health metrics, improving pipeline reliability from 87% to 96% uptime.
- Implemented **Tableau** and **Looker** dashboards with real-time alerting mechanisms using **Grafana** and **Datadog** to monitor data freshness and quality, enabling the business to detect issues within 15 minutes instead of days.

- Drove adoption of dbt for data transformation logic, establishing version control and CI/CD integration with **GitHub Actions**, reducing deployment time for analytics changes from 2 hours to under 10 minutes.
- Mentored junior engineers on analytics engineering best practices, including star schema design, data modeling for event and transactional data, and working with healthcare data governance requirements like HIPAA compliance.
- Partnered with data engineering teams to transition v1 analytics data sets to long-term ownership, creating comprehensive documentation and automated tests using **Pytest** to ensure data quality and reduce maintenance burden.
- Supported external reporting initiatives by building and validating extraction pipelines for healthcare data exports, ensuring compliance with data privacy standards and successful delivery to third-party analytics platforms.
- Reduced time to insights by 35% through optimization of **BigQuery** and **Snowflake** queries, implementing materialized views and incremental loading patterns that scaled to process 500M+ daily healthcare events.
- Established data quality checks and automated alerting for upstream data set changes, coordinating with business teams to manage schema migrations and claims logic rewrites without disrupting analytics workflows.
- Built reusable **Python** and **Pandas** utilities for common data transformation patterns, reducing code duplication across the analytics team and accelerating new data asset development by 25%.
- Collaborated across fintech and healthcare domains to design extensible data models that anticipate future business needs, eliminating the need for major rework when requirements evolved mid-project.

CareSimple

Oct 2022–May 2024

Senior Full Stack Engineer

- Spearheaded the creation of population health analytics data sets from raw remote patient monitoring (RPM) and hospital-at-home data, using **SQL** and **Python** to build reliable assets serving clinical and administrative stakeholders.
- Designed and maintained data models using star schema principles to support chronic care management analytics, translating domain expertise from clinical teams into computational logic that reduced reporting errors by 40%.
- Built **Airflow** orchestration layers to schedule and automate ETL workflows for healthcare data integration, processing daily streams of RPM sensor data, EHR events, and claims records across multiple payer systems.
- Implemented data governance and quality frameworks using custom **Python** and **Pandas** validation scripts, catching data anomalies before they reached downstream analytics and reducing invalid records by 28%.
- Developed **Looker** blocks and **Tableau** dashboards to surface population health metrics for hospital systems and physician groups, enabling clinicians to identify high-risk patients and improve intervention strategies.

- Collaborated with product and business teams to capture analytics requirements for new virtual care workflows, translating feature requests into data logic that supported immediate business value while remaining extensible for future use cases.
- Optimized **Snowflake** queries for large-scale healthcare claims analysis, reducing query execution time from 8 minutes to 1.2 minutes through partitioning, caching strategies, and materialized views.
- Established alerting mechanisms using **Grafana** and **Datadog** to monitor healthcare data pipeline health, ensuring 99.2% on-time delivery of critical data sets serving payer and provider decision-making.
- Documented data lineage and asset relationships using internal data cataloging tools, enabling 35+ analysts to confidently explore the analytics landscape without redundant queries or computational overhead.
- Supported external compliance reporting for healthcare organizations by building validated extraction pipelines for quality metrics, readmission rates, and cost data, ensuring audit-trail integrity for regulatory submissions.

Shopee

Jul 2018–Sep 2022

Full Stack Engineer

- Played a key role in building the marketplace analytics infrastructure for a 295-million-user e-commerce platform, developing ETL pipelines in **Python** and **PySpark** that ingested seller transaction, customer behavior, and logistics data into a centralized **PostgreSQL** and later **BigQuery** warehouse.
- Designed and implemented star schema data models for transactional and event data, enabling 200+ analysts and product managers to explore user behavior, seller performance, and supply chain metrics without creating duplicate logic or analytical silos.
- Built **Airflow** DAGs to orchestrate daily and real-time data loading jobs, processing 80 million+ transactions per day across eight countries and managing data dependencies across 30+ upstream systems with 94% on-time completion rates.
- Collaborated with marketplace teams to translate business requirements into data logic for seller segmentation, customer cohort analysis, and supply chain optimization, delivering analytics assets that informed decisions reducing marketing spend by 12% while maintaining engagement.
- Optimized data transformation pipelines by migrating batch jobs from **Spark** on Hadoop to **BigQuery** native transformations, reducing processing time by 55% and infrastructure costs by 18% while improving data freshness from daily to near-real-time.
- Implemented **Tableau** and internal dashboarding solutions to surface key performance indicators for the fintech and logistics segments, enabling finance and operations teams to monitor fraud rates, payment delays, and delivery SLAs in real-time.
- Developed monitoring and alerting systems using **Datadog** to track data pipeline health, detecting anomalies in transaction volumes, seller uploads, and inventory syncs that could indicate upstream issues, reducing mean time to detection by 60%.
- Created reusable **Python** and **Pandas** libraries for common transformation patterns used across seller analytics, buyer segmentation, and logistics tracking, standardizing data logic and reducing on-boarding time for new analytics use cases.

- Supported cross-functional alignment on the single source of truth for seller metrics by partnering with finance and operations teams to codify business logic for revenue attribution, commission calculations, and performance rankings.
- Managed data quality and schema evolution when upstream systems changed, coordinating with infrastructure teams to ensure backwards compatibility and communicating impact to downstream analytics consumers without disrupting workflows.

Altitude Games

Nov 2014–Jul 2018

Full Stack Engineer

- Built backend services and data pipelines for free-to-play mobile games with 5 million+ global users, using **Python**, **Java**, and **SQL** to collect and process game event data, user progression metrics, and in-app purchase transactions.
- Designed and maintained **PostgreSQL** schemas to store player behavioral data, in-game economy metrics, and monetization events, establishing clear data models that supported analytics queries for retention, engagement, and revenue optimization.
- Implemented automated data export pipelines using **Python** and **Pandas** to enable non-technical game designers to analyze A/B test results and user cohort performance without writing complex queries.
- Developed reporting dashboards using custom **JavaScript** and **HTML/CSS** visualizations that surfaced daily active users, monetization funnels, and level difficulty metrics, enabling rapid iteration on game balance and economy.
- Collaborated with game design and monetization teams to translate business questions into data requirements, supporting experiments that improved player retention by 18% through informed difficulty scaling and progression pacing.
- Optimized query performance by implementing caching layers using **Redis** for frequently accessed aggregations, reducing dashboard load times from 12 seconds to 2 seconds and improving analyst productivity.
- Built versioning and change management into game event schemas, coordinating with backend teams to ensure that new game features and mechanics generated properly structured event data that analytics could consume without rework.
- Maintained data integrity through automated validation checks, catching issues where event payloads were malformed or missing critical fields before they reached analytics, reducing data-driven decision latency.

EDUCATION

Ateneo de Manila University

Aug 2010–Jul 2014

Bachelor of Science in Computer Science